

# Taints & Tolerations

**Taints:** Node-level "repellents" that repel pods unless they tolerate the taint.

**Tolerations:** Pod-level "acceptance" rules that allow pods to be scheduled on tainted nodes.

## Step 1: Apply Taints to Worker Nodes

These taints will prevent pods from scheduling unless they tolerate the taint.

```
controlplane ~ → kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
controlplane	Ready	control-plane	70m	v1.33.0
node01	Ready	<none>	70m	v1.33.0
node02	Ready	<none>	69m	v1.33.0

```
controlplane ~ → kubectl taint nodes node01 gpu=true:NoSchedule
kubectl taint nodes node02 gpu=false:NoSchedule
node/node01 tainted
node/node02 tainted
```

## Step 2: Deploy a Pod Without Toleration

```
controlplane ~ → vi nginx-tainted.yaml
```

```
controlplane ~ → cat nginx-tainted.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-tainted
spec:
  containers:
  - name: nginx
    image: nginx
```

```
controlplane ~ → kubectl apply -f nginx-tainted.yaml
kubectl get pods
kubectl describe pod nginx-tainted
pod/nginx-tainted created
```

Type	Reason	Age	From	Message
----	-----	----	----	-----
Warning	FailedScheduling	1s	default-scheduler	0/3 nodes are available: 1 node(s) had untolerated taint {gpu: false}, 1 node(s) had untolerated taint {gpu: true}, 1 node(s) had untolerated taint {node-role.kubernetes.io/control-plane: }. preemption: 0/3 nodes are available: 3 Preemption is not helpful for scheduling.

```
controlplane ~ → cat nginx-tainted.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-tainted
spec:
  tolerations:
  - key: "gpu"
    operator: "Equal"
    value: "true"
    effect: "NoSchedule"
  containers:
  - name: nginx
    image: nginx
```

```
controlplane ~ → kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	REA
nginx-tainted	1/1	Running	0	36s	172.17.1.2	node01	<none>		<no

### Step 3: Deploy Redis Pod to Control Plane

```
controlplane ~ → cat redis-on-controlplane.yaml
apiVersion: v1
kind: Pod
metadata:
  name: redis-control
spec:
  containers:
  - name: redis
    image: redis
```

```
controlplane ~ ➔ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE     NOMINATED NODE   REA
DINESS GATES
nginx-tainted                        1/1     Running   0           18m   172.17.1.2    node01   <none>           <no
ne>
redis-control                        0/1     Pending   0           65s   <none>        <none>   <none>           <no
ne>
```